

Credit_Risk_Modeling

May 26, 2018

0.0.1 Invoke Packages

```
In [148]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import GridSearchCV
from mlxtend.plotting import plot_confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import confusion_matrix, mean_squared_error, classification_report
%matplotlib inline
```

0.0.2 Reading dataframe and cleaning the noise

```
In [149]: loans_2007 = pd.read_csv(r"./databank/loans_2007.csv", low_memory = False)

# Removing rows with duplicate data
loans_2007 = loans_2007.drop_duplicates()
# Display first row of the dataframe
print(loans_2007.info())
# Display dataframe size
print("\nTotal number of features : ", loans_2007.shape[1])
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 42538 entries, 0 to 42537
Data columns (total 52 columns):
id                42538 non-null object
member_id         42535 non-null float64
loan_amnt         42535 non-null float64
funded_amnt       42535 non-null float64
funded_amnt_inv   42535 non-null float64
term              42535 non-null object
int_rate          42535 non-null object
installment       42535 non-null float64
grade             42535 non-null object
sub_grade         42535 non-null object
emp_title         39909 non-null object
```

emp_length	41423	non-null	object
home_ownership	42535	non-null	object
annual_inc	42531	non-null	float64
verification_status	42535	non-null	object
issue_d	42535	non-null	object
loan_status	42535	non-null	object
pymnt_plan	42535	non-null	object
purpose	42535	non-null	object
title	42522	non-null	object
zip_code	42535	non-null	object
addr_state	42535	non-null	object
dti	42535	non-null	float64
delinq_2yrs	42506	non-null	float64
earliest_cr_line	42506	non-null	object
inq_last_6mths	42506	non-null	float64
open_acc	42506	non-null	float64
pub_rec	42506	non-null	float64
revol_bal	42535	non-null	float64
revol_util	42445	non-null	object
total_acc	42506	non-null	float64
initial_list_status	42535	non-null	object
out_prncp	42535	non-null	float64
out_prncp_inv	42535	non-null	float64
total_pymnt	42535	non-null	float64
total_pymnt_inv	42535	non-null	float64
total_rec_prncp	42535	non-null	float64
total_rec_int	42535	non-null	float64
total_rec_late_fee	42535	non-null	float64
recoveries	42535	non-null	float64
collection_recovery_fee	42535	non-null	float64
last_pymnt_d	42452	non-null	object
last_pymnt_amnt	42535	non-null	float64
last_credit_pull_d	42531	non-null	object
collections_12_mths_ex_med	42390	non-null	float64
policy_code	42535	non-null	float64
application_type	42535	non-null	object
acc_now_delinq	42506	non-null	float64
chargeoff_within_12_mths	42390	non-null	float64
delinq_amnt	42506	non-null	float64
pub_rec_bankruptcies	41170	non-null	float64
tax_liens	42430	non-null	float64

dtypes: float64(30), object(22)
memory usage: 17.2+ MB
None

Total number of features : 52

0.0.3 Cleaning dataframe by columns carrying irrelevant information

```
In [150]: cols = ["id", "member_id", "funded_amnt", "funded_amnt_inv", "grade", "sub_grade", "issue_d", "zip_code", "out_prncp", "out_prncp_inv", "total_pymnt", "total_pymnt_inv", "total_rec_prncp", "total_rec_int", "total_rec_late_fee", "recoveries", "col", "last_pymnt_d", "last_pymnt_amnt"]
loans_2007 = loans_2007.drop(cols, axis=1)
loans_2007.shape[1]
```

Out[150]: 32

0.0.4 Selecting Target Column

```
In [151]: # Selecting "loan_status" as a target column
loans_2007["loan_status"].value_counts()
```

```
Out[151]: Fully Paid                                33136
Charged Off                                         5634
Does not meet the credit policy. Status:Fully Paid  1988
Current                                             961
Does not meet the credit policy. Status:Charged Off  761
Late (31-120 days)                                24
In Grace Period                                    20
Late (16-30 days)                                  8
Default                                             3
Name: loan_status, dtype: int64
```

0.0.5 Cleaning the target column

```
In [152]: # Only first two values above in "Loan_Status" are found to be useful for our modeling
keep = ["Fully Paid", "Charged Off"]
map_dict = {"Fully Paid": 1, "Charged Off": 0}

loans_2007 = loans_2007[(loans_2007["loan_status"] == keep[0]) | (loans_2007["loan_status"] == keep[1])]

loans_2007["loan_status"] = loans_2007["loan_status"].map(map_dict)
loans_2007["loan_status"].unique()
```

Out[152]: array([1, 0], dtype=int64)

0.0.6 Removing the columns with single value

```
In [153]: # Since single valued columns carry no useful information for credibility evaluation
drop_columns = []

for col in loans_2007.columns:
    current = loans_2007[col].dropna()
    if len(current.unique()) == 1:
        drop_columns.append(col)
```

```
loans_2007 = loans_2007.drop(drop_columns, axis=1)
loans_2007.columns
```

```
Out[153]: Index(['loan_amnt', 'term', 'int_rate', 'installment', 'emp_length',
                'home_ownership', 'annual_inc', 'verification_status', 'loan_status',
                'purpose', 'title', 'addr_state', 'dti', 'delinq_2yrs',
                'earliest_cr_line', 'inq_last_6mths', 'open_acc', 'pub_rec',
                'revol_bal', 'revol_util', 'total_acc', 'last_credit_pull_d',
                'pub_rec_bankruptcies'],
                dtype='object')
```

0.0.7 Checking the dataframe for null values

```
In [154]: # Display list of columns carrying null values (in descending order)
null_counts = loans_2007.isnull().sum().sort_values(ascending=False)
null_counts[0:10]
```

```
Out[154]: emp_length      1036
pub_rec_bankruptcies     697
revol_util                50
title                    11
last_credit_pull_d        2
purpose                   0
term                      0
int_rate                  0
installment               0
home_ownership            0
dtype: int64
```

0.0.8 Verifying the null values proportion

```
In [155]: # Show columns with number of null values as percentage of total training examples
# (in descending order)
null_percent = (loans_2007.isnull().sum()/loans_2007.shape[0]).sort_values(ascending=False)
null_percent[0:5]
```

```
Out[155]: emp_length      2.672169
pub_rec_bankruptcies     1.797782
revol_util               0.128966
title                   0.028372
last_credit_pull_d       0.005159
dtype: float64
```

Consider clearing columns with more less 1% NULL value while removing the columns with above 1% NULL values.

0.0.9 Further Cleaning/Removing the dataframe

```
In [156]: loans_2007 = loans_2007.drop('pub_rec_bankruptcies', axis=1)
          loans_2007 = loans_2007.dropna(subset=['title', 'revol_util', 'last_credit_pull_d'],
                                          axis=0, how='any')
          print("Total NaN values in the dataframe: ", loans_2007.isnull().sum().sum())
          print('pub_rec_bankruptcies' in loans_2007)
```

```
Total NaN values in the dataframe: 1032
False
```

0.0.10 Review the columns that carry object type data

```
In [157]: object_columns_df = loans_2007.select_dtypes(include=["object"])
          object_columns_df.iloc[0]
```

```
Out[157]: term                36 months
          int_rate              10.65%
          emp_length            10+ years
          home_ownership        RENT
          verification_status    Verified
          purpose                credit_card
          title                  Computer
          addr_state             AZ
          earliest_cr_line       Jan-1985
          revol_util              83.7%
          last_credit_pull_d      Jun-2016
          Name: 0, dtype: object
```

The following columns seem to carry categorical data that could be helpful for our regression analysis later:

- term
- emp_length
- home_ownership
- verification_status
- purpose
- title

The above results show that the following columns carry type of information that we can convert into numerical type for further analysis:

- int_rate
- revol_util

The remaining table seem to be not much useful and should be drop for our data frame. They are:

- addr_state
- earliest_cr_line
- last_credit_pull_d

0.0.11 Verify the columns with categorical data

```
In [158]: cols = ['term', 'emp_length', 'home_ownership', 'verification_status', 'purpose', 't
```

```
    for i in cols:
        print("\nUnique Counts for " + i + " :\n", object_columns_df[i].value_counts())
```

Unique Counts for term :

36 months	29040
60 months	9667

Name: term, dtype: int64

Unique Counts for emp_length :

10+ years	8545
< 1 year	4513
2 years	4303
3 years	4022
4 years	3353
5 years	3202
1 year	3176
6 years	2177
7 years	1714
8 years	1442
9 years	1228

Name: emp_length, dtype: int64

Unique Counts for home_ownership :

RENT	18513
MORTGAGE	17111
OWN	2984
OTHER	96
NONE	3

Name: home_ownership, dtype: int64

Unique Counts for verification_status :

Not Verified	16696
Verified	12289
Source Verified	9722

Name: verification_status, dtype: int64

Unique Counts for purpose :

debt_consolidation	18130
credit_card	5039
other	3864
home_improvement	2897
major_purchase	2154
small_business	1762
car	1510

wedding	929
medical	680
moving	576
vacation	375
house	369
educational	320
renewable_energy	102

Name: purpose, dtype: int64

Unique Counts for title :

Debt Consolidation	2104
Debt Consolidation Loan	1632
Personal Loan	642
Consolidation	494
debt consolidation	485
Credit Card Consolidation	353
Home Improvement	346
Debt consolidation	324
Small Business Loan	310
Credit Card Loan	305
Personal	302
Consolidation Loan	251
Home Improvement Loan	234
personal loan	227
personal	211
Loan	208
Wedding Loan	201
Car Loan	195
consolidation	193
Other Loan	181
Credit Card Payoff	150
Wedding	149
Credit Card Refinance	143
Major Purchase Loan	139
Consolidate	125
Medical	118
Credit Card	115
home improvement	107
My Loan	92
Credit Cards	91
...	
money honey	1
Working capital business loan	1
Graduation/Travel Expenses	1
Home Improvements	1
Making ends meet	1
amexloan	1
HudsonKK	1

New kitchen	1
Small business loan	1
COSMETIC	1
Dept Consolidation Loan Success	1
Freash Start!	1
Triathlete	1
Hearing Aids	1
business plan debt consultation	1
Condo Renovation Loan	1
pay off high interest debt	1
CREATING FINACIAL INDEPENDENCE FOR MY CHILDREN AND ME	1
Home sale buyout	1
KDC_Loan123	1
Car loan and pay off my family	1
Credit Card Relieve	1
Family Support	1
Catrin's wish	1
Furniture	1
Monty	1
Discover payoff	1
New Job Relocate Debt Until IL Home Sell	1
missvi1	1
The consolidator	1

Name: title, Length: 19331, dtype: int64

It appears that column 'purpose' and 'title' carry overlapping information. However, considering the duplication of loan categories under the 'title' column, let's keep the 'purpose' column and drop the 'title' one.

Column 'emp_length' carries multiple category data.

0.0.12 Drop less useful columns with duplicate/overlapping data

```
In [159]: cols = ['addr_state', 'earliest_cr_line', 'last_credit_pull_d', 'title']
         loans_2007 = loans_2007.drop(cols, axis=1)
```

0.0.13 Assign data variables to columns

```
In [160]: cat_columns = ['home_ownership', 'verification_status', 'purpose', 'term']
         loans_2007_dummies = pd.get_dummies(loans_2007[cat_columns])
         loans_2007 = pd.concat([loans_2007, loans_2007_dummies], axis=1)
```

```
# Removinng the original non-dummy columns
loans_2007 = loans_2007.drop(cat_columns, axis=1)
loans_2007.columns
```

```
Out[160]: Index(['loan_amnt', 'int_rate', 'installment', 'emp_length', 'annual_inc',
               'loan_status', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'open_acc',
               'pub_rec', 'revol_bal', 'revol_util', 'total_acc',
```



```

'home_ownership_MORTGAGE', 'home_ownership_NONE',
'home_ownership_OTHER', 'home_ownership_OWN', 'home_ownership_RENT',
'verification_status_Not Verified',
'verification_status_Source Verified', 'verification_status_Verified',
'purpose_car', 'purpose_credit_card', 'purpose_debt_consolidation',
'purpose_educational', 'purpose_home_improvement', 'purpose_house',
'purpose_major_purchase', 'purpose_medical', 'purpose_moving',
'purpose_other', 'purpose_renewable_energy', 'purpose_small_business',
'purpose_vacation', 'purpose_wedding', 'term_ 36 months',
'term_ 60 months'],
dtype='object')

```

0.0.14 Mapping column with multiple category data

```

In [161]: mapping_dict = {
    "emp_length": {
        "10+ years": 10,
        "9 years": 9,
        "8 years": 8,
        "7 years": 7,
        "6 years": 6,
        "5 years": 5,
        "4 years": 4,
        "3 years": 3,
        "2 years": 2,
        "1 year": 1,
        "< 1 year": 0,
        np.nan: 0
    }
}

loans_2007 = loans_2007.replace(mapping_dict)
np.sort(loans_2007['emp_length'].unique().astype('int64'))

Out[161]: array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10], dtype=int64)

```

0.0.15 Convert Columns String type Numeric Data to Real Numbers

```

In [162]: loans_2007['int_rate'] = loans_2007['int_rate'].str.rstrip('%').astype('float')
    loans_2007['revol_util'] = loans_2007['revol_util'].str.rstrip('%').astype('float')

```

0.0.16 Preserving completely cleansed dataframe in .csv format

```

In [163]: loans_2007.to_csv("cleaned_loans_2007.csv", index=False)
    loans_2007.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 38707 entries, 0 to 39785
Data columns (total 38 columns):

```

loan_amnt	38707 non-null float64
int_rate	38707 non-null float64
installment	38707 non-null float64
emp_length	38707 non-null int64
annual_inc	38707 non-null float64
loan_status	38707 non-null int64
dti	38707 non-null float64
delinq_2yrs	38707 non-null float64
inq_last_6mths	38707 non-null float64
open_acc	38707 non-null float64
pub_rec	38707 non-null float64
revol_bal	38707 non-null float64
revol_util	38707 non-null float64
total_acc	38707 non-null float64
home_ownership_MORTGAGE	38707 non-null uint8
home_ownership_NONE	38707 non-null uint8
home_ownership_OTHER	38707 non-null uint8
home_ownership_OWN	38707 non-null uint8
home_ownership_RENT	38707 non-null uint8
verification_status_Not Verified	38707 non-null uint8
verification_status_Source Verified	38707 non-null uint8
verification_status_Verified	38707 non-null uint8
purpose_car	38707 non-null uint8
purpose_credit_card	38707 non-null uint8
purpose_debt_consolidation	38707 non-null uint8
purpose_educational	38707 non-null uint8
purpose_home_improvement	38707 non-null uint8
purpose_house	38707 non-null uint8
purpose_major_purchase	38707 non-null uint8
purpose_medical	38707 non-null uint8
purpose_moving	38707 non-null uint8
purpose_other	38707 non-null uint8
purpose_renewable_energy	38707 non-null uint8
purpose_small_business	38707 non-null uint8
purpose_vacation	38707 non-null uint8
purpose_wedding	38707 non-null uint8
term_ 36 months	38707 non-null uint8
term_ 60 months	38707 non-null uint8

dtypes: float64(12), int64(2), uint8(24)

memory usage: 5.3 MB

0.0.17 Finalizing the feature columns

```
In [164]: features = loans_2007.columns.values
```


0.0.18 Evaluating the data skewness

```
In [165]: target_vals = loans_2007["loan_status"].value_counts()
          print(target_vals)

          print("The binary data ratio = 1 :", target_vals[1] / target_vals[0])
          sns.countplot(x="loan_status", data=loans_2007)
```

```
1    33092
0     5615
Name: loan_status, dtype: int64
The binary data ratio = 1 : 5.893499554764025
```

```
Out[165]: <matplotlib.axes._subplots.AxesSubplot at 0xce180b8>
```



This clearly indicates that our data is skewed.

Let's consider a logistic regression model that predict "1" as loan would be pay-off on time and "0" as loan would not be pay-off on time.

Considering the conservative investors approach:

- * We must emphasize an algorithm that can avoid predicting "False Positive" which can suggest them any risky investment as a safe option. This may, otherwise, cause severe business losses.

- * Similarly, it is acceptable that the algorithm is little lenient with predicting "False Negative" so that some safe investments could end up with wrong rejections. This may cause slip of some safe business opportunity.

0.0.19 Develop Logistic Regression Model and Fit Verification

```
In [166]: # split train/test/cv as 60/20/20
train_set = loans_2007.iloc[:23224]
test_set = loans_2007.iloc[23225:30966]
cross_val = loans_2007.iloc[30967:]

# Find the logistic regression model using sklearn
logistic_model = LogisticRegression()
logistic_model.fit(train_set[features], train_set["loan_status"])
predictions_test = logistic_model.predict(test_set[features])
mse_vals_test = mean_squared_error(predictions_test, test_set["loan_status"])

mse_vals_test
```

Out[166]: 0.1331869267536494

0.0.20 Verifying the Logistic Reg Model over CV Set

```
In [167]: # Find the MSE for CV

logistic_model = LogisticRegression()
logistic_model.fit(train_set[features], train_set["loan_status"])
predictions_cv = logistic_model.predict(cross_val[features])
mse_vals_cv = mean_squared_error(predictions_cv, cross_val["loan_status"])

mse_vals_cv
```

Out[167]: 0.13217054263565892

0.0.21 Checking the Model Generalization using Another Classifier

```
In [168]: # Check for model generalization using scikit-learn package and carry out K-Fold
# Cross Validation

logistic_model = LogisticRegression()
predictions_tot = pd.Series(cross_val_predict(logistic_model, X=loans_2007[features],
                                              y=loans_2007["loan_status"], cv=11))
mse_vals_tot = mean_squared_error(predictions_tot, loans_2007["loan_status"])

mse_vals_tot
```

Out[168]: 0.14573591340067688

The error results for model predictability over test set and cross-validation set using two different classifier. This shows a better model generalization for predictability over outside data.

0.0.22 Check Prediction Integrity using Evaluation Matrix

In [169]: *# Calc for TP/FP/TN/FN*

For Logistic Reg Classifier over Test Set

```
tp = len(predictions_test[(predictions_test == 1) & (test_set["loan_status"] == 1)])
tn = len(predictions_test[(predictions_test == 0) & (test_set["loan_status"] == 0)])
fp = len(predictions_test[(predictions_test == 1) & (test_set["loan_status"] == 0)])
fn = len(predictions_test[(predictions_test == 0) & (test_set["loan_status"] == 1)])

precision_test = tp / (tp + fp)
recall_test = tp / (tp + fn)
accuracy_test = (tp + tn) / len(test_set['loan_status'])
fpr_test = fp / (fp + tn)
tpr_test = tp / (tp + fn)
F1_score_test = 2 * precision_test * recall_test / (precision_test + recall_test)

print("Precision over Test Set:", precision_test)
print("Recall over Test Set:", recall_test)
print("Accuracy over Test Set:", accuracy_test)
print("FPR over Test Set:", fpr_test)
print("TPR over Test Set:", tpr_test)
print("F1 Score over Test Set:", F1_score_test)

# Verify the calc using sklearn function
plot_confusion_matrix(confusion_matrix(test_set["loan_status"], predictions_test))
print(classification_report(test_set['loan_status'], predictions_test))
```

Precision over Test Set: 0.8683084899546338

Recall over Test Set: 0.9977658623771224


Accuracy over Test Set: 0.8668130732463506

FPR over Test Set: 0.9892891918208374

TPR over Test Set: 0.9977658623771224

F1 Score over Test Set: 0.9285466768313811

	precision	recall	f1-score	support
0	0.42	0.01	0.02	1027
1	0.87	1.00	0.93	6714
avg / total	0.81	0.87	0.81	7741



```
In [170]: # Calc for TP/FP/TN/FN
```

```
# For Logistic Reg Classifier over Cross Validation Set
```

```
tp = len(predictions_cv[(predictions_cv == 1) & (cross_val["loan_status"] == 1)])
tn = len(predictions_cv[(predictions_cv == 0) & (cross_val["loan_status"] == 0)])
fp = len(predictions_cv[(predictions_cv == 1) & (cross_val["loan_status"] == 0)])
fn = len(predictions_cv[(predictions_cv == 0) & (cross_val["loan_status"] == 1)])
```

```
precision_cv = tp / (tp + fp)
recall_cv = tp / (tp + fn)
accuracy_cv = (tp + tn) / len(cross_val['loan_status'])
fpr_cv = fp / (fp + tn)
tpr_cv = tp / (tp + fn)
F1_score_cv = 2 * precision_cv * recall_cv / (precision_cv + recall_cv)
```

```
print("Precision over Cross Validation Set:", precision_cv)
print("Recall over Cross Validation Set:", recall_cv)
print("Accuracy over Cross Validation Set:", accuracy_cv)
print("FPR over Cross Validation Set:", fpr_cv)
print("TPR over Cross Validation Set:", tpr_cv)
print("F1 Score over Cross Validation Set:", F1_score_cv)
```

```
# Verify the calc using sklearn function
```

```
plot_confusion_matrix(confusion_matrix(cross_val["loan_status"], predictions_cv))
print(classification_report(cross_val['loan_status'], predictions_cv))
```

```
Precision over Cross Validation Set: 0.8683563946721842
```


```
Recall over Cross Validation Set: 0.9992559523809523
```

```
Accuracy over Cross Validation Set: 0.8678294573643411
```

```
FPR over Cross Validation Set: 0.9980392156862745
```

TPR over Cross Validation Set: 0.9992559523809523
 F1 Score over Cross Validation Set: 0.9292188472981389

	precision	recall	f1-score	support
0	0.29	0.00	0.00	1020
1	0.87	1.00	0.93	6720
avg / total	0.79	0.87	0.81	7740



In [171]: # Calc for TP/FP/TN/FN

For K-Fold Cross Validation Classifier over entire data set

```

tp = len(predictions_tot[(predictions_tot == 1) & (loans_2007["loan_status"] == 1)])
tn = len(predictions_tot[(predictions_tot == 0) & (loans_2007["loan_status"] == 0)])
fp = len(predictions_tot[(predictions_tot == 1) & (loans_2007["loan_status"] == 0)])
fn = len(predictions_tot[(predictions_tot == 0) & (loans_2007["loan_status"] == 1)])

precision_tot = tp / (tp + fp)
recall_tot = tp / (tp + fn)
accuracy_tot = (tp + tn) / len(loans_2007['loan_status'])
fpr_tot = fp / (fp + tn)
tpr_tot = tp / (tp + fn)
F1_score_tot = 2 * precision_tot * recall_tot / (precision_tot + recall_tot)

print("Precision by K-Fold Classify:", precision_tot)
print("Recall by K-Fold Classify:", recall_tot)
print("Accuracy by K-Fold Classify:", accuracy_tot)
print("FPR by K-Fold Classify:", fpr_tot)

```

```

print("TPR by K-Fold Classify:", tpr_tot)
print("F1 Score by K-Fold Classify:", F1_score_tot)

# Verify the calc using sklearn function
plot_confusion_matrix(confusion_matrix(loans_2007["loan_status"], predictions_tot))
print(classification_report(loans_2007['loan_status'], predictions_tot))

```

Precision by K-Fold Classify: 0.8556253660614451

Recall by K-Fold Classify: 0.9982295937383526


Accuracy by K-Fold Classify: 0.8306766217996745

FPR by K-Fold Classify: 0.9974250505793636

TPR by K-Fold Classify: 0.9982295937383526

F1 Score by K-Fold Classify: 0.9214427019123255

	precision	recall	f1-score	support
0	0.32	0.00	0.01	5615
1	0.86	1.00	0.92	33092
avg / total	0.78	0.85	0.79	38707



0.0.23 Results

```

In [172]: output_cols = ["Test Set", "Cross Validation Set", "K-Fold Classify"]
          output_index = ["precision", "recall", "accuracy", "fpr", "tpr", "F1_score", "MSE"]
          output = pd.DataFrame(index = output_index, columns = output_cols)

          output['Test Set'] = [precision_test, recall_test, accuracy_test, fpr_test, tpr_test,
                               F1_score_test, mse_vals_test]
          output['Cross Validation Set'] = [precision_cv, recall_cv, accuracy_cv, fpr_cv, tpr_cv,

```



```

                                F1_score_cv, mse_vals_cv]
output['K-Fold Classify'] = [precision_tot, recall_tot, accuracy_tot, fpr_tot, tpr_tot,
                                F1_score_tot, mse_vals_tot]
output

```

```

Out[172]:
      Test Set  Cross Validation Set  K-Fold Classify
precision  0.868308                0.868356        0.855625
recall    0.997766                0.999256        0.998230
accuracy  0.866813                0.867829        0.830677
fpr       0.989289                0.998039        0.997425
tpr       0.997766                0.999256        0.998230
F1_score  0.928547                0.929219        0.921443
MSE       0.133187                0.132171        0.145736

```

0.0.24 Results Review

* The precision of 0.5 indicates that only 50

* The prediction algorithm seems to be biased towards the positive outcome. This also gets reflected in the high recall value. This could not gain confident to make business/investment decision by the conservative investors.

* High value of FPR shows the positive bias of the predictive algorithm. The 99

0.0.25 Revise Model

```

In [173]: # Need to reduce "False Positive" to improve predictive reliability
          # To enhance conservative predictive outcome, let's set-up penalty for wrongly label
          penalty = {0:6, 1:1}
          # Set up Logistic regression model
          logistic_model = LogisticRegression(class_weight = penalty)
          logistic_model.fit(train_set[features], train_set["loan_status"])
          predictions_test_rev = logistic_model.predict(test_set[features])
          mse_vals_test_rev = mean_squared_error(predictions_test_rev, test_set["loan_status"])

          mse_vals_test_rev

```

```

Out[173]: 0.0

```

0.0.26 Re-verify the Logistic Reg Revised Model

```

In [174]: # Find the MSE for CV

          logistic_model.fit(train_set[features], train_set["loan_status"])
          predictions_cv_rev = logistic_model.predict(cross_val[features])
          mse_vals_cv_rev = mean_squared_error(predictions_cv_rev, cross_val["loan_status"])

          mse_vals_cv_rev

```

```

Out[174]: 0.0

```

0.0.27 Re-Check (Manually) the Revised Model Generalization using Class Weights Classifier

In order to accomplish splitting the skewed data appropriately into partitions, every class should be applied with different weight so that all partitions are eventually left out for testing. We have data skewed in the order of 1:5.89. It should be noted that RandomForestClassifier suites to allow for class weighting.

We should manage the class weights manually to find a better balance between the "False Positives" and "False Negatives". The F1 score is a good indicator of tradeoff between the precision and recall.

```
In [175]: # Check for model generalization using scikit-learn GridSearchCV classifier

wts = np.linspace(0.05, 0.95, 20)

# Initialize the classifier (presumming 3-Folds will provide a reasonable balance be
# the accuracy and computation time)
grdcv = GridSearchCV(
    estimator=LogisticRegression(),
    param_grid={
        'class_weight': [{0: x, 1: 1.0-x} for x in wts]
    },
    scoring='f1',
    cv=3)


grd_rslt = grdcv.fit(loans_2007[features], loans_2007['loan_status'])
print("Best parameters : %s" % grd_rslt.best_params_)

# Plot the weights vs F1_score
df_wt = pd.DataFrame({ 'score': grd_rslt.cv_results_['mean_test_score'],
                        'weight': wts })
df_wt.plot(x='weight')

# Initializing the Logistic Regression Cross Validation
logistic_model_rev = LogisticRegression(**grd_rslt.best_params_)
logistic_model_rev.fit(train_set[features], train_set['loan_status'])
predictions_tot_rev = logistic_model_rev.predict(cross_val[features])
mse_vals_tot_rev = mean_squared_error(predictions_tot_rev, cross_val["loan_status"])

print("MSE for GridSearchCV Cross-Validation Check =", mse_vals_tot_rev)

Best parameters : {'class_weight': {0: 0.7605263157894736, 1: 0.2394736842105264}}
MSE for GridSearchCV Cross-Validation Check = 0.0
```



0.0.28 Re-Check Prediction Integrity using Evaluation Matrix

In [176]: *# Calc for TP/FP/TN/FN*

For Logistic Reg Classifier over Test Set

```
tp = len(predictions_test_rev[(predictions_test_rev == 1) & (test_set["loan_status"] == 1)])
tn = len(predictions_test_rev[(predictions_test_rev == 0) & (test_set["loan_status"] == 0)])
fp = len(predictions_test_rev[(predictions_test_rev == 1) & (test_set["loan_status"] == 0)])
fn = len(predictions_test_rev[(predictions_test_rev == 0) & (test_set["loan_status"] == 1)])
```

```
precision_test_rev = tp / (tp + fp)
```

```
recall_test_rev = tp / (tp + fn)
```

```
accuracy_test_rev = (tp + tn) / len(test_set['loan_status'])
```

```
fpr_test_rev = fp / (fp + tn)
```

```
tpr_test_rev = tp / (tp + fn)
```

```
F1_score_test_rev = 2 * precision_test_rev * recall_test_rev / (precision_test_rev + recall_test_rev)
```

```
print("Precision over Test Set:", precision_test_rev)
```

```
print("Recall over Test Set:", recall_test_rev)
```

```
print("Accuracy over Test Set:", accuracy_test_rev)
```

```
print("FPR over Test Set:", fpr_test_rev)
```

```
print("TPR over Test Set:", tpr_test_rev)
```

```
print("F1 Score over Test Set:", F1_score_test_rev)
```

```
# Verify the calc using sklearn function
plot_confusion_matrix(confusion_matrix(test_set["loan_status"], predictions_test_rev),
print(classification_report(test_set['loan_status'], predictions_test_rev))
```

Precision over Test Set: 1.0

Recall over Test Set: 1.0


Accuracy over Test Set: 1.0

FPR over Test Set: 0.0

TPR over Test Set: 1.0

F1 Score over Test Set: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1027
1	1.00	1.00	1.00	6714
avg / total	1.00	1.00	1.00	7741



In [177]: *# Calc for TP/FP/TN/FN*

For Logistic Reg Classifier over Cross Validation Set

```
tp = len(predictions_cv_rev[(predictions_cv_rev == 1) & (cross_val["loan_status"] == 1)])
tn = len(predictions_cv_rev[(predictions_cv_rev == 0) & (cross_val["loan_status"] == 0)])
fp = len(predictions_cv_rev[(predictions_cv_rev == 1) & (cross_val["loan_status"] == 0)])
fn = len(predictions_cv_rev[(predictions_cv_rev == 0) & (cross_val["loan_status"] == 1)])
```

```
precision_cv_rev = tp / (tp + fp)
```

```
recall_cv_rev = tp / (tp + fn)
```

```
accuracy_cv_rev = (tp + tn) / len(cross_val['loan_status'])
```

```

fpr_cv_rev = fp / (fp + tn)
tpr_cv_rev = tp / (tp + fn)
F1_score_cv_rev = 2 * precision_cv_rev * recall_cv_rev / (precision_cv_rev + recall_cv_rev)

print("Precision over Cross Validation Set:", precision_cv_rev)
print("Recall over Cross Validation Set:", recall_cv_rev)
print("Accuracy over Cross Validation Set:", accuracy_cv_rev)
print("FPR over Cross Validation Set:", fpr_cv_rev)
print("TPR over Cross Validation Set:", tpr_cv_rev)
print("F1 Score over Cross Validation Set:", F1_score_cv_rev)

# Verify the calc using sklearn function
plot_confusion_matrix(confusion_matrix(cross_val["loan_status"], predictions_cv_rev))
print(classification_report(cross_val['loan_status'], predictions_cv_rev))


```

```

Precision over Cross Validation Set: 1.0
Recall over Cross Validation Set: 1.0
Accuracy over Cross Validation Set: 1.0
FPR over Cross Validation Set: 0.0
TPR over Cross Validation Set: 1.0
F1 Score over Cross Validation Set: 1.0

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1020
1	1.00	1.00	1.00	6720
avg / total	1.00	1.00	1.00	7740



```

In [178]: # Calc for TP/FP/TN/FN
          # For GridSearchCV Classifier with Logistic Reg over the cross validation data set

```

```

tp = len(predictions_tot_rev[(predictions_tot_rev == 1) & (cross_val["loan_status"] =
tn = len(predictions_tot_rev[(predictions_tot_rev == 0) & (cross_val["loan_status"] =
fp = len(predictions_tot_rev[(predictions_tot_rev == 1) & (cross_val["loan_status"] =
fn = len(predictions_tot_rev[(predictions_tot_rev == 0) & (cross_val["loan_status"] =

precision_tot_rev = tp / (tp + fp)
recall_tot_rev = tp / (tp + fn)
accuracy_tot_rev = (tp + tn) / len(cross_val['loan_status'])
fpr_tot_rev = fp / (fp + tn)
tpr_tot_rev = tp / (tp + fn)
F1_score_tot_rev = 2 * precision_tot_rev * recall_tot_rev / (precision_tot_rev +
                                                              recall_tot_rev)

print("Precision by GridSearchCV Classifier:", precision_tot_rev)
print("Recall by GridSearchCV Classifier:", recall_tot_rev)
print("Accuracy by GridSearchCVClassifier:", accuracy_tot_rev)
print("FPR by GridSearchCV Classifier:", fpr_tot_rev)
print("TPR by GridSearchCV Classifier:", tpr_tot_rev)
print("F1 Score by GridSearchCV Classifier:", F1_score_tot_rev)

# Verify the calc using sklearn function
plot_confusion_matrix(confusion_matrix(cross_val["loan_status"], predictions_tot_rev),
print(classification_report(cross_val['loan_status'], predictions_tot_rev))


```

```

Precision by GridSearchCV Classifier: 1.0
Recall by GridSearchCV Classifier: 1.0
Accuracy by GridSearchCVClassifier: 1.0
FPR by GridSearchCV Classifier: 0.0
TPR by GridSearchCV Classifier: 1.0
F1 Score by GridSearchCV Classifier: 1.0

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1020
1	1.00	1.00	1.00	6720
avg / total	1.00	1.00	1.00	7740



0.0.29 Revised Model Results Summary

```
In [179]: output_cols = ["Test Set", "Cross Validation Set", "GridSearch CV"]
          output_rev = pd.DataFrame(index = output_index, columns = output_cols)

          output_rev['Test Set'] = [precision_test_rev, recall_test_rev, accuracy_test_rev, fpr_test_rev, tpr_test_rev, F1_score_test_rev, mse_vals_test_rev]
          output_rev['Cross Validation Set'] = [precision_cv_rev, recall_cv_rev, accuracy_cv_rev, fpr_cv_rev, tpr_cv_rev, F1_score_cv_rev, mse_vals_cv_rev]
          output_rev["GridSearch CV"] = [precision_tot_rev, recall_tot_rev, accuracy_tot_rev, fpr_tot_rev, tpr_tot_rev, F1_score_tot_rev, mse_vals_tot_rev]

          output_rev
```

```
Out[179]:
```

	Test Set	Cross Validation Set	GridSearch CV
precision	1.0	1.0	1.0
recall	1.0	1.0	1.0
accuracy	1.0	1.0	1.0
fpr	0.0	0.0	0.0
tpr	1.0	1.0	1.0
F1_score	1.0	1.0	1.0
MSE	0.0	0.0	0.0

0.0.30 Conclusion

* The results above show generalization of the revised model over the unknown data set.

* Precision value of 1.0 indicates that this model predicts ONLY actual positive examples as positive outcome.

* Recall value of 1.0 is an indication that this model captures ALL positive examples correctly.

* F1-Score is an indicative of perfect 'Precision' and 'Recall' scores.

* The zero value of fpr confirms that the model is absolutely not predicting "False Positive". This confirms the reliability of the algorithm and its alignment with our need for conservative business proposition to avoid any risky investment option. The checks for model predictive behavior above assure this model would recommend only loan options that are safe and will be paid off on time.

* The TPR value of 1.0 is an indicative that this model will show ALL of the safe loan opportunities that will be paid off on time. This will give an utmost confidence to the conservative investors and encourage them to treat any/all of the recommended business opportunity as risk free to do business.